

HW 4 Solutions

ENGR 21, Fall 2026.

October 1, 2026

Solutions

1 Python function to read IEEE floats

Computers store numbers in memory using a series of zeros and ones. These zeros and ones are then interpreted as floating-point numbers according to the standard scheme that we learned about in class.

In this question, you will write a Python function that can interpret a string of zeros and ones as a decimal number using the **IEEE standard**. Your function should work for 16-bit, 32-bit and 64-bit numbers.

Tip

Name your function `readIEEEfloat`. It should take a single argument in the form of a string and should return the value as a number.

Solution

Convert binary integer to decimal integer

```
def bin_to_dec_int(num):
    # Converts binary integer to decimal integer
    # num is a string, interpreted as a binary integer.
    s = len(num)
    total_value = 0
    for index in range(s):
        power = s-index-1
        digit = int(num[index])
        value = (2 ** power) * digit
        total_value += value
    return total_value
```

Convert binary fraction to decimal fraction

```
def bin_to_dec_frac(num):
    # Converts binary fraction to decimal fraction.
    # num is a string of 0's and 1's.
    # If num = 1010, this function interprets it as
    #  $1 * (1/2)^1 + 0 * (1/2)^2 + 1 * (1/2)^3 + 0 * (1/2)^4$ .
    # The first bit multiplies (1/2). The second bit multiplies (1/4). The third bit multiplies (1/8).
    # for example, in the IEEE format, a 16-bit number could have the significand
    # 1.0011010101
    # This function can be used to read the string of bits AFTER the "decimal point" and
    # returns the value of the resulting fraction in decimal form as a 'float'.
    # The returned value must always be less than 1.
    s = len(num)
    total_value = 0
    for index in range(s):
        power = -index - 1
        digit = int(num[index])
        value = (2 ** power) * digit
        total_value += value
    return total_value
```

Combine

```
def readIEEEfloat(num):
    # Interpret num, type string, as a 16-bit, 32-bit or 64-bit binary number stored
    # in IEEE format. Return the answer as an object of type float.

    if len(num) == 16 or len(num) == 32 or len(num) == 64:
        # First, take out the sign bit
        signbit = int(num[0])
        sign = -1 if signbit == 1 else 1

        # Then, take out the exponent and significand bits
        if len(num) == 16:
            exponent = bin_to_dec_int(num[1:6])
            significand = bin_to_dec_frac(num[6:])
            bias = 15
        elif len(num) == 32:
            exponent = bin_to_dec_int(num[1:9])
            significand = bin_to_dec_frac(num[9:])
            bias = 127
        elif len(num) == 64:
            exponent = bin_to_dec_int(num[1:12])
            significand = bin_to_dec_frac(num[12:])
            bias = 1023

        # Then put it together.
        return sign * (1+significand) * 2 ** (exponent - bias)
    else:
        print("Input must be a string of length 16, 32 or 64")
        return None
```

2 Floating-point numbers

2.1 Conversion to fractional form

Interpret the following 16-bit floating-point numbers as irreducible fractions. Also write an approximate decimal representation of your fractions.

1. 0b1001110000101110, a.k.a. 0x1C2E
2. 0b0100110001101010, a.k.a. 0x4C6A

Solutions

1. We first divide up the 16 bits of 0b1001110000101110 into:
 - Sign bit: 1 — this number is negative
 - Exponent bits: 00111 — in decimal form, this is 7.

- Significand bits: 0000101110

Using this information, we can write this number in the following way:

$$-1.0000101110_2 \times 2^{7-15}$$

which is equal to

$$-(1 + 2^{-5} + 2^{-7} + 2^{-8} + 2^{-9}) \times 2^{-8}$$

We can now write it out as fractions:

$$\begin{aligned} & -\left(1 + \frac{1}{32} + \frac{1}{128} + \frac{1}{256} + \frac{1}{512}\right) \times \frac{1}{256} \\ &= -\left(\frac{512}{512} + \frac{16}{512} + \frac{4}{512} + \frac{2}{512} + \frac{1}{512}\right) \times \frac{1}{256} \\ &= -\left(\frac{535}{512}\right) \times \frac{1}{256} = \boxed{-\frac{535}{131072}} \approx -0.00408 \end{aligned}$$

2. We first divide up the 16 bits of 0b0100110001101010 into:

- Sign bit: 0 — this number is positive
- Exponent bits: 10011 — in decimal form, this is 19.
- Significand bits: 0001101010

Using this information, we can write this number in the following way:

$$+1.0001101010_2 \times 2^{19-15}$$

which is equal to

$$(1 + 2^{-4} + 2^{-5} + 2^{-7} + 2^{-9}) \times 2^4$$

We can now write it out as fractions:

$$\begin{aligned} & \left(1 + \frac{1}{16} + \frac{1}{32} + \frac{1}{128} + \frac{1}{512}\right) \times 16 \\ &= \left(\frac{512}{512} + \frac{32}{512} + \frac{16}{512} + \frac{4}{512} + \frac{1}{512}\right) \times 16 \\ &= \left(\frac{565}{512}\right) \times 16 = \boxed{\frac{565}{32}} \approx 17.656 \end{aligned}$$

2.2 Large and small Numbers

Use the web app [Float Toy](#) to fill the following table. All numbers should be positive, and should be given in scientific notation in decimal form up to the number of significant figures given by the web app or up to 3 significant figures after the decimal point, whichever is smaller.

 Solutions

Number	16-bit	32-bit	64-bit
Smallest subnormal non-zero number	5.97×10^{-8}	1×10^{-45}	5×10^{-324}
Second smallest subnormal non-zero number	1.2×10^{-7}	2.8×10^{-45}	1×10^{-323}
Smallest normal* non-zero number	6.104×10^{-5}	1.175×10^{-38}	2.225×10^{-308}
Second smallest normal* non-zero number	6.11×10^{-5}	1.175×10^{-38}	2.225×10^{-308}
Largest non-NaN number	6.550×10^4	3.403×10^{38}	1.797×10^{308}
Second largest non-NaN number	6.547×10^4	3.403×10^{38}	1.797×10^{308}

* i.e., not subnormal.

3 Gap Size in floating point numbers

3.1 Gaps between 16-bit floating-point binary numbers

We would like to tabulate the ‘gap size’, which can also be called the increment size, between floating-point numbers. Two of the values in this table have been filled from the relevant sections of Lecture 8 [here](#) and [here](#).

Between	and	the gap is	or equivalently
2^{-14}	2^{-13}	...	
2^{-13}	2^{-12}	...	
2^{-12}	2^{-11}	...	
2^{-11}	2^{-10}	...	
2^{-10}	2^{-9}	...	
2^{-9}	2^{-8}	...	
2^{-8}	2^{-7}	...	
2^{-7}	2^{-6}	...	
2^{-6}	2^{-5}	...	
2^{-5}	2^{-4}	...	
2^{-4}	2^{-3}	...	
2^{-3}	2^{-2}	...	
2^{-2}	2^{-1}	...	
2^{-1}	2^0	...	
2^0	2^1	...	
2^1	2^2	...	
2^2	2^3	...	
2^3	2^4	2^{-7}	7.812×10^{-3}
2^4	2^5	...	
2^5	2^6	...	
2^6	2^7	...	
2^7	2^8	...	
2^8	2^9	...	
2^9	2^{10}	...	
2^{10}	2^{11}	...	
2^{11}	2^{12}	2^1	2
2^{12}	2^{13}	...	
2^{13}	2^{14}	...	
2^{14}	2^{15}	...	
2^{15}	2^{16}	...	

 Solutions

Between	and	the gap is	or equivalently
2^{-14}	2^{-13}	2^{-24}	3.051×10^{-5}
2^{-13}	2^{-12}	2^{-23}	1.192×10^{-7}
2^{-12}	2^{-11}	2^{-22}	2.384×10^{-7}
2^{-11}	2^{-10}	2^{-21}	4.768×10^{-7}
2^{-10}	2^{-9}	2^{-20}	9.536×10^{-7}
2^{-9}	2^{-8}	2^{-19}	1.907×10^{-6}
2^{-8}	2^{-7}	2^{-18}	3.814×10^{-6}
2^{-7}	2^{-6}	2^{-17}	7.629×10^{-6}
2^{-6}	2^{-5}	2^{-16}	1.525×10^{-5}
2^{-5}	2^{-4}	2^{-15}	3.051×10^{-5}
2^{-4}	2^{-3}	2^{-14}	6.103×10^{-5}
2^{-3}	2^{-2}	2^{-13}	1.220×10^{-4}
2^{-2}	2^{-1}	2^{-12}	2.441×10^{-4}
2^{-1}	2^0	2^{-11}	4.882×10^{-4}
2^0	2^1	2^{-10}	9.765×10^{-4}
2^1	2^2	2^{-9}	1.953×10^{-3}
2^2	2^3	2^{-8}	3.906×10^{-3}
2^3	2^4	2^{-7}	7.812×10^{-3}
2^4	2^5	2^{-6}	1.562×10^{-2}
2^5	2^6	2^{-5}	3.125×10^{-2}
2^6	2^7	2^{-4}	6.250×10^{-2}
2^7	2^8	2^{-3}	0.125
2^8	2^9	2^{-2}	0.25
2^9	2^{10}	2^{-1}	0.5
2^{10}	2^{11}	2^0	1
2^{11}	2^{12}	2^1	2
2^{12}	2^{13}	2^2	4
2^{13}	2^{14}	2^3	8
2^{14}	2^{15}	2^4	16
2^{15}	2^{16}	2^5	32

3.2 Precision for floats

The $\pm$ symbol is often used to indicate the precision of a known quantity. For floating point numbers, it is appropriate to use $x \pm y$ to represent a number, where y is half of the gap size in that part of the number line.

For example, we saw in lecture 8 that between 8 and 16, 16-bit floating point numbers have a gap size of $1/128$ or 0.0078125. Half of this number is 0.00390625. If we now consider the 16-bit floating point number given by 0100100000001100, its value can be found to be equal to $259/32 = 8.09375$. However, it would not be correct to write this number in decimal form with 5 significant figures after the decimal point. To see why, let us write it in $\pm$ notation as follows.

$$8.09375 \pm 0.00390625$$

Thus, the 16-bit float 0100100000001100 would be used by a computer to represent any number in the range from $8.09375 - 0.00390625$ to $8.09375 + 0.00390625$ shown above. For example,

- 8.09672
- 8.09585
- 8.09175
- 8.09079

are all within the range shown above. These numbers agree in the first two digits after the decimal point, but in subsequent places, they disagree. Therefore, when representing the floating-point number 0100100000001100 in decimal form, it would not be appropriate to write more significant figures after the second one after the decimal point. Thus, the correct way to express it in decimal form would be 8.09, not 8.09375.

Using the discussion above as a template, write the following floating-point numbers in fractional form and then in decimal form using the appropriate number of significant figures after the decimal point.

1. The 16-bit float given by 0b0001110000010000

Solution

Interpreting this number, we find that it is

$$1.0000010000_2 \times 2^{7-15} = \left(1 + \frac{1}{2^6}\right) \times 2^{-8}$$

$$\frac{65}{64} \times \frac{1}{256} = \frac{65}{16384}$$

In decimal form, this can be written as 0.00396728515625000. The gap size in this part of the number line, for 16-bit floats, is $2^{-18} \approx 0.000003814697$. Half of this number would be 0.00000190734.

We now compare the top and bottom of the range of possible values that this floating-point number is supposed to represent. These are: 0.00396919249 and 0.00396537781 respectively. These two agree up to the fifth place after the decimal point, but no further. Therefore, we should write this number as

$$0.00396 \quad \text{or} \quad 3.96 \times 10^{-3}.$$

Alternatively, you could also choose to round this up to 3.97×10^{-3} because the sixth place after the decimal point is always 5 or greater.

2. The 16-bit float given by 0x3C10

 Solution

Interpreting this number, we find that it is

$$\left(1 + \frac{1}{2^6}\right) \times 2^0$$
$$\frac{65}{64} \times 1 = \frac{8388609}{536870912}$$

In decimal form, this can be written as 1.015625000. The gap size in this part of the number line, for 16-bit floats, is $2^{-10} \approx 0.0009765625$. Half of this number would be 0.00048828125. We now compare the top and bottom of the range of possible values that this floating-point number is supposed to represent. These are: 1.01611328125 and 1.0151468 respectively. These two agree up to the second place after the decimal point, but no further. Therefore, we should write this number as

1.01

and with no further significant figures.

4 The Farmer Was Replaced

Download the save file [here](#) to complete this part of the homework.

4.1 Zero location

Write a function called `zeroloc()` that takes the drone back to the bottom-left corner.

 Solution

```
def zeroloc():
    xdist = get_pos_x()
    ydist = get_pos_y()
    for i in range(xdist):
        move(West)
    for i in range(ydist):
        move(South)
    return None
```

4.2 Random plants

Write a script that plants a random plant (out of five options: Grass, Tree, Bush, Carrot, and Pumpkin) on each spot of the 6x6 farm. It should only do this once, without repeating.

Your script should `import` your `zeroloc` function from a different file, i.e., the first line of your script should be something like `from <file name> import zeroloc`.

The second line of your script should call the `zeroloc` function.

Solution

```
from positioning import zeroloc
zeroloc()

crops = {
    "p": Entities.Pumpkin,
    "t": Entities.Tree,
    "g": Entities.Grass,
    "c": Entities.Carrot,
    "b": Entities.Bush
}

possible_crops = ["p", "t", "b", "c", "g"]

for row in range(get_world_size()):
    for col in range(get_world_size()):
        if get_ground_type() == Grounds.Grassland:
            till()
            if can_harvest():
                harvest()
            ind = random() * len(possible_crops) // 1
            plant(crops[possible_crops[ind]])
            move(East)
            move(North)
```

4.3 Waiting for crops to grow

Write a function called `plant_and_wait` that takes as argument a plant ‘object’, such as `Entities.Carrot`.

When this function is called with argument `Entities.Tree`, its behavior should be as shown in the gif below, i.e., it should plant a tree and do flips until the tree has fully grown, and then harvest it.

When this function is used to plant pumpkins, its behavior needs to be different. It should wait to check if the pumpkin grows up to be rotten, and if so, it should start over. As an illustration of this behavior, the following gif shows what your drone should do if you run the following code:

```
while True:
    plant_and_wait(Entities.Pumpkin)
```

 Warning

This function should not move the drone. It **should** till the soil if it is grassland (remember that tilling grass turns it into soil, but tilling soil turns it back into grass!)

 Solution

```
def plant_and_wait(crop):
    if get_ground_type() == Grounds.Grassland:
        till()
    if not crop == Entities.Pumpkin:
        if can_harvest():
            harvest()
        plant(crop)
        while not can_harvest():
            do_a_flip()
        harvest()
    else:
        plant(crop)
        while not can_harvest():
            do_a_flip()
            if get_entity_type() == Entities.Dead_Pumpkin:
                print("Death")
                plant(Entities.Pumpkin)
        harvest()
```